



PennState

CMPSC 311 - Introduction to Systems Programming

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Introduction to Concurrency

Add WeChat edu_assist_pro

Suman Saha

(Slides are mostly by Professors Patrick McDaniel and Abutalib Aghayev)



Sequential Programming



- Processing a network connection as it arrives and fulfilling the exchange completely is sequential processing
 - i.e., connections are processed in sequence of arrival

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
listen_fd = Listen(  
while(1){  
    client_fd = accep  
    buf = read(client_fd);  
    write(client_fd, buf);  
    close(client_fd);  
}
```

Add WeChat edu_assist_pro

Whither sequential?



- Benefits

- simple to implement
- very little persistent state to maintain
- few complex error conditions

- Disadvantages

- Sometimes poor performance
 - one slow client causes all others to block
 - poor utilization of network, CPU

Think about it this way: if the class took the final exam sequentially, it would take 25 days to complete

An alternate design ...



PennState

- Why not process multiple requests at the same time, interleaving processing while waiting for other actions (e.g., read requests) to complete?
- This is known as concurrent processing ...
 - Process multiple requests
- Approaches to concurrent server design
 - Asynchronous servers (`select()`)
 - Multiprocess servers (`fork()`)
 - Multithreaded servers (`threads`)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Concurrency with processes



PennState

- The server process blocks on `accept()`, waiting for a new client to connect
 - when a new connection arrives, the parent calls `fork()` to create another process
 - the child process handles that new connection, and `exit()`'s when the co
- Children become “zom”
 - `wait()` to “reap” children

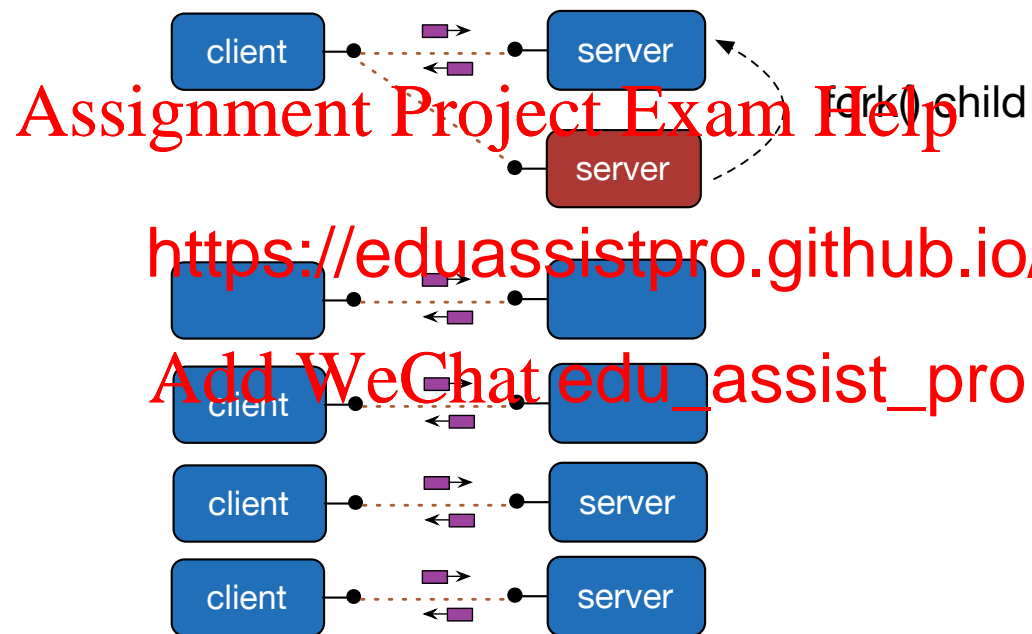
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Graphically



fork()



PennState

- The fork function creates a new process by duplicating the calling process.

- The new **child** process is an exact duplicate of the calling **parent** process, except that it has its own **queue**.
- The `fork()` function returns
 - 0 (zero) for the child process.
 - The child's process ID in the parent code.

Idea: think about duplicating the process state and running

Process control



PennState

- Parent
 - fork (pid == child PID)
 - wait for child to complete (maybe)
- Child
 - begins at fork (pid == 0)
 - runs until done
 - calls exit

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

exit()



- The `exit` causes normal process termination with a return value

```
void exit(int status);
```

- Where

- status is sent to the to

- Note: `exit` vs. `return` from a C program

- `return` is a language keyword
 - returns (possibly a value) from a function
 - `exit` is a function, eventually calls `_exit` a system call
 - terminates process immediately, returns status to parent

- `exit` and `return` are similar in main function

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro



wait()



PennState

- The `wait` function is used to wait for state changes in a child of the calling process (e.g., to terminate)

Assignment Project Exam Help

```
pid_t wait(int *status);
```

- Where

- returns the process ID
- `status` is return value set by the child pro

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Putting it all together ...



PennState

```
int main(void)
{
    printf("starting parent process\n");
    pid_t pid = fork();

    if (pid < 0) {
        perror("fork");
        exit(EXIT_FAILURE);
    }
    if (pid == 0) {
        printf("child processing\n");
        sleep(50);
        exit(19);
    } else {
        printf("parent forked a child with pid = %d\n", pid);
        int status;
        wait(&status);
        printf("child exited with status = %d\n", WEXITSTATUS(status));
    }
    return 0;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Concurrency with processes




- Benefits

- almost as simple as sequential
- in fact, most of the code is identical!
- parallel execution; good
- often better security (isolation)

- Disadvantages

- processes are heavyweight
- relatively slow to fork
- context switching latency is high
- communication between processes is complicated



Concurrency with threads



PennState

- A single process handles all of the connections
 - ... but ... a **parent thread** forks (dispatches) a new thread to handle each connection
 - the **child thread**:
 - handles the new co
 - exits when the conne

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Note: you can create as many threads as you want (up to a system limit)

Threads



PennState

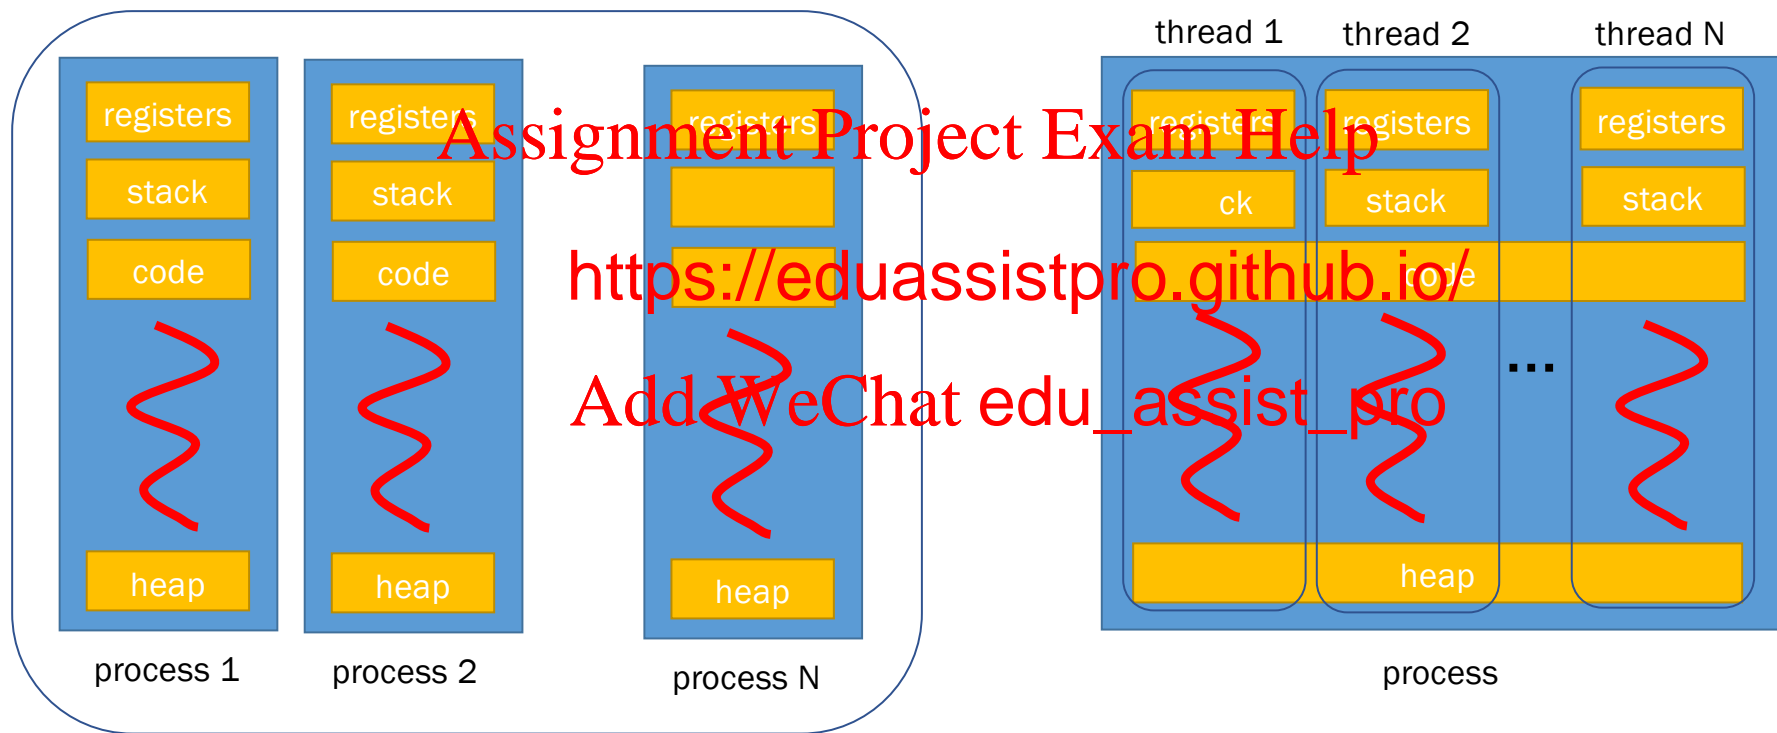
- A thread is defined as an independent stream of instructions that can be scheduled to run as such by the operating system.
 - To the software developer, the concept of a "procedure" that runs independently from its main program.
 - To go one step further, imagine a <https://eduassistpro.github.io/> procedures. Now imagine all of these procedures to be scheduled to run simultaneously and/or independently by the system. That would describe a "*multi-threaded*" program.

Idea: “forking” multiple threads of execution in one process!

Multiple Processes vs Multiple Threads



PennState



Multiple processes

Threads (cont.)



PennState

Assignment Project Exam Help

<https://eduassistpro.github.io/>

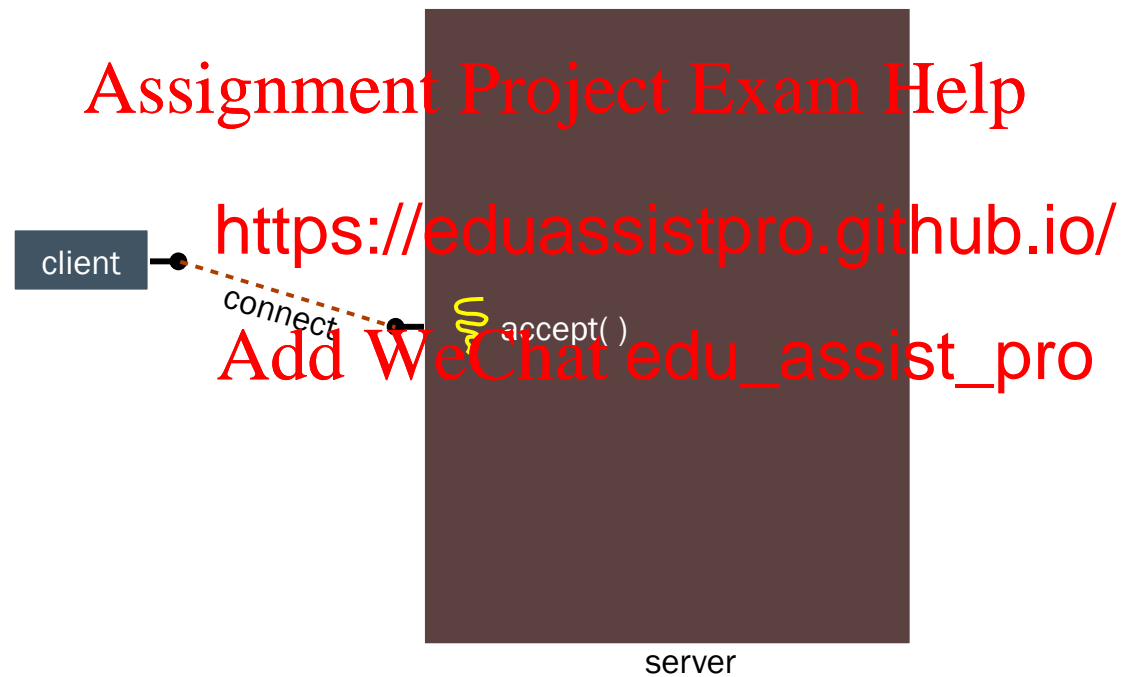
Add WeChat  `accept()` edu_assist_pro

server

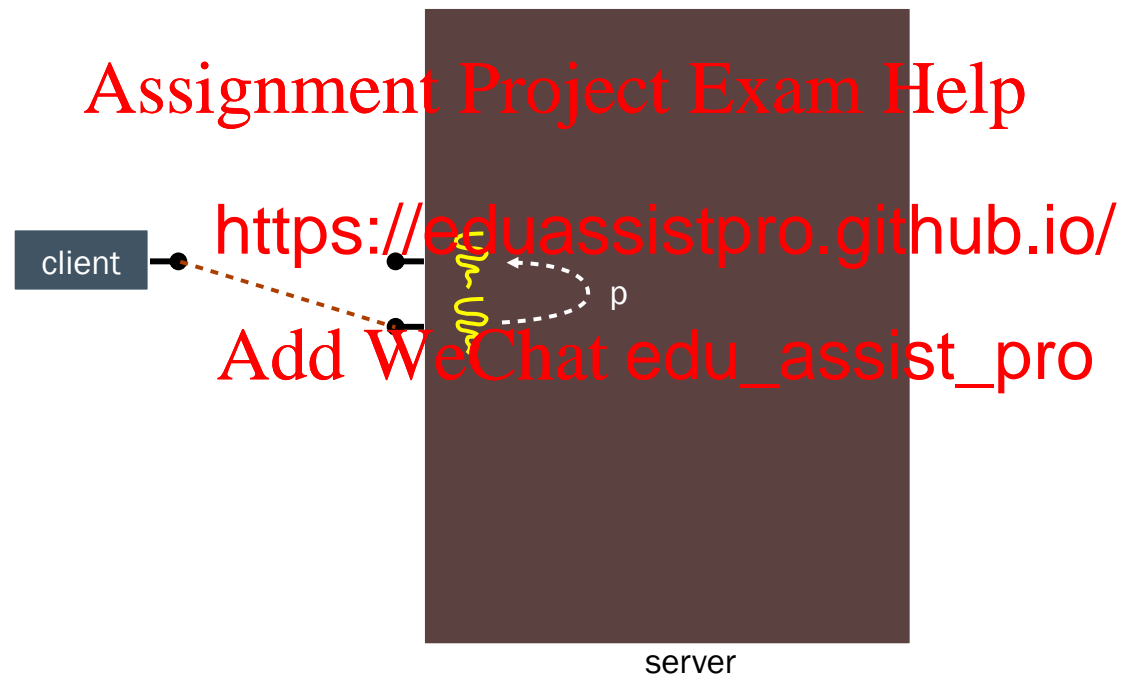
Threads (cont.)



PennState



Threads (cont.)



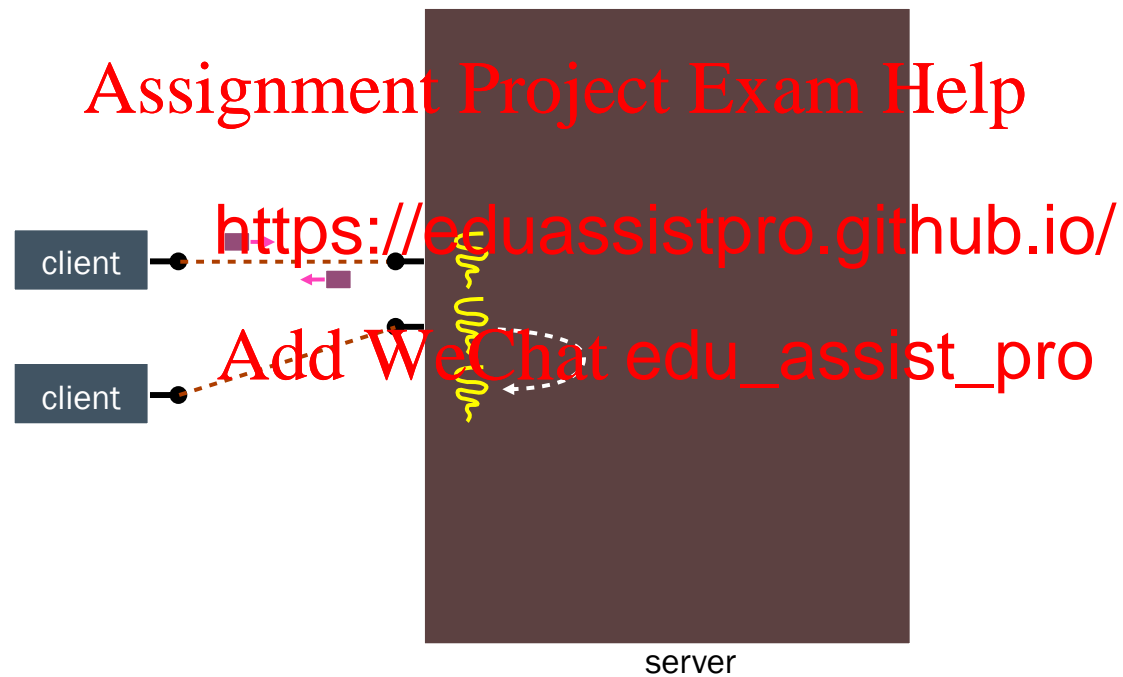
Threads (cont.)



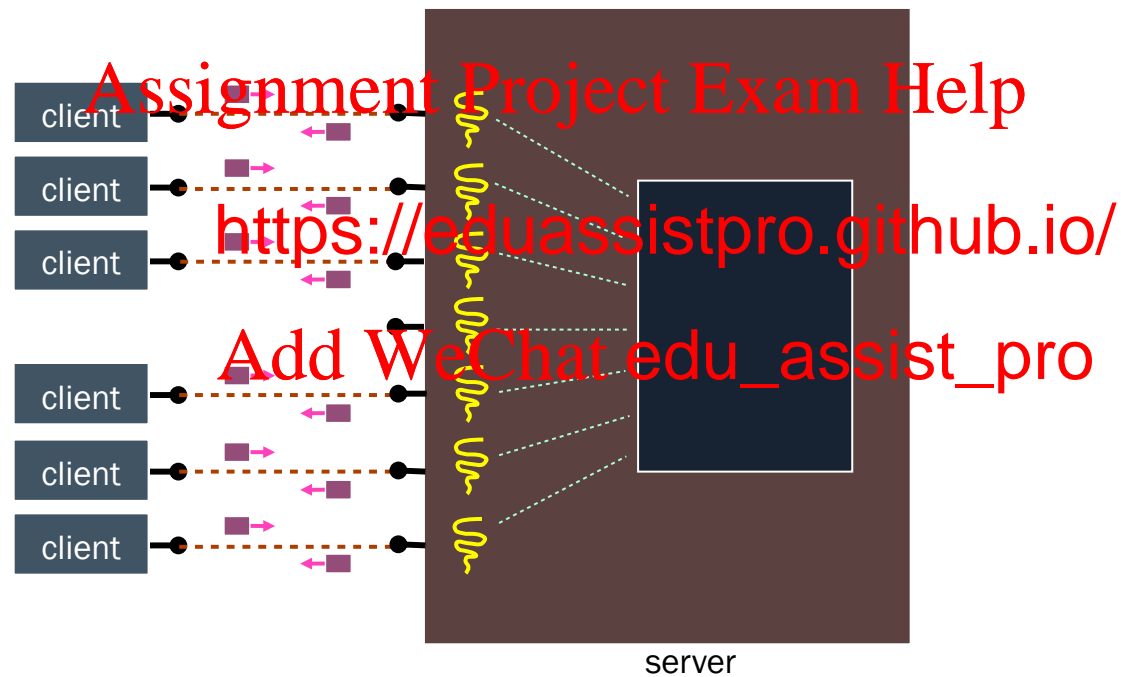
Assignment Project Exam Help



Threads (cont.)



Threads (cont.)



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

UNIX Process

... and with threads

Threads (cont.)



PennState

- This independent flow of control is accomplished because a thread maintains its own:

- Stack pointer
- Registers
- Scheduling properties (
- Set of pending and blocked signals
- Thread specific data.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Thread Summary



- Exists within a process and uses the process resources
- Has its own independent flow of control as long as its parent process exists and the OS supports it
- Duplicates only the essential resources and runs independently
“schedulable”
- May share the process resources with other threads that act equally independently (and dependently)
- Dies if the parent process dies - or something similar
- Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.

Caveats



- Because threads within the same process share resources:
 - Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads.
 - Two pointers having the same value
 - Reading and writing to the same memory location, and therefore requires explicit synchronization by the programmer.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Warning: shared data between threads can cause conflicts, deadlocks, etc.

Thread control



PennState

- main
 - `pthread_create()` (create thread)
 - wait for thread to finish via `pthread_join()` (maybe)
- thread
 - begins at function point
 - runs until the `return` or `pthread_exit`
- Library support
 - `#include <pthread.h>`
 - Compile with option `-lpthread` to link with the pthread library

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

pthread_create()



PennState

- The `pthread_create` function starts a new thread in the calling process.

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  e (void *),
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Where,

- `thread` is a pthread library structure hold
- `attr` is a set of attributes to apply to the thread
- `start_routine` is the thread function pointer
- `arg` is an opaque data pointer to pass to thread

Add WeChat edu_assist_pro

Thread with no arguments



PennState

```
void *func(void *arg) {
    printf("Hello from thread %lx\n", pthread_self());
    return NULL;
}

int main(void) {
    pthread_t t1, t2, t3;
    printf("main thread %lx start\n", pthread_self());
    pthread_create(&t1, NULL, func, NULL);
    pthread_create(&t2, NULL, func, NULL);
    pthread_create(&t3, NULL, func, NULL);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    pthread_join(t3, NULL);
    return 0;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- Always check return values (omitted for brevity)
- Thread becomes alive in `pthread_create` – may even run before it returns

Thread with one argument



PennState

```
void *func(void *arg) {  
    char *s = (char *) arg;  
    printf("Hello from thread %s\n", s);  
    return NULL;  
}  
  
int main(void) {  
    pthread_t t1, t2, t3;  
    pthread_create(&t1, NULL,  
    pthread_create(&t2, NULL,  
    pthread_create(&t3, NULL, func, "b");  
    pthread_join(t1, NULL);  
    pthread_join(t2, NULL);  
    pthread_join(t3, NULL);  
    return 0;  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- Run the above program in a loop to observe indeterminate scheduling

Thread with multiple arguments



PennState

```
typedef struct {
    int num;
    const char *str;
} foo_t;

void *func(void *arg) {
    foo_t *val = arg;
    printf("thread %lx was passed v          val->str);
    return NULL;
}

int main(void) {
    foo_t v = {5678, "bar"};
    pthread_t t;
    printf("main thread %lx starting a new thread\n", pthread_self());
    pthread_create(&t, NULL, func, &v);
    pthread_join(t, NULL);
    return 0;
}
```

```
$ ./t3
main thread 7f46f18d0740 starting a new thread
thead 7f46f18cf700 was passed values 5678, bar
[0$ euclid:/tmp/master]
$
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- The above is effectively a procedure call – real programs are more complex

pthread_join()



- The `pthread_join` function waits for the thread specified by `thread` to terminate.

```
int pthread_join(pthread_t thread, void **retval);
```

- Where,

- `thread` is a pthread_t
- `retval` is a double pointer return value

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Returning values from a thread



PennState

```
typedef struct {
    int num;
    char *str;
} foo_t;
```

```
void *func(void *arg) {
    foo_t *a = arg;
    foo_t *b = malloc(sizeof(foo_t));
    b->num = a->num * 2;
    b->str = malloc(strlen(a->str) + 1);
    strcpy(b->str, a->str);
    for (char *p = b->str; *p; ++p)
        *p = toupper(*p);
    return b;
}
```

```
int main(void) {
    foo_t v = {5678, "bar"}, *p;
    pthread_t t;
    printf("main thread %lx starting a new thread\n", pthread_self());
    pthread_create(&t, NULL, func, &v);
    pthread_join(t, (void **) &p);
    printf("thread returned num = %d, str = %s\n", p->num, p->str);
    return 0;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Returning values from a thread



PennState

```
typedef struct {  
    int num;  
    const char *str;  
} foo_t;
```

```
void *func(void *arg) {  
    foo_t p;  
    // fill p  
    return &p;  
}
```

```
int main(void) {  
    foo_t v = {5678, "bar"}, *p;  
    pthread_t t;  
    printf("main thread %lx starting a new thread\n", (void *)0);  
    pthread_create(&t, NULL, func, &v);  
    pthread_join(t, (void **) &p);  
    printf("thread returned num = %d, str = %s\n", p->num, p->str);  
    return 0;  
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- The above will segfault! Do not return a pointer to a stack-allocated variable

pthread_exit()



- The `pthread_exit` function terminates the calling thread and returns a value


Assignment Project Exam Help

`void pthread_exit(void *retval);`

- Where,
 - `retval` is a pointer to

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Threads accessing shared data



PennState

```
static int counter = 0;

void *func(void *arg) {
    for (int i = 0; i < 5000; ++i)
        ++counter;
    return NULL;
}

int main(void)
{
    pthread_t t1, t2;
    printf("counter = %d\n", counter);
    pthread_create(&t1, NULL, func, NULL);
    pthread_create(&t2, NULL, func, NULL);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("both threads completed, counter = %d\n", counter);
    return 0;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- What will this program output?

What is happening? A race condition!



PennState

- **Race condition** happens when the outcome of a program depends on the interleaving of the execution of multiple threads accessing **critical section**
- **Critical section** is a piece of code that accesses a shared variable and must not be concurrently executed

```
mov    0x2e50(%rip),%eax    # 4014 <counter>
add     $0x1,%eax
mov     %eax,0x2e47(%rip)
```

- Each instruction executed atomically
- Multiple threads executing the above instructions can result in different interleavings (and outcomes) due to uncontrollable OS scheduling

One Possible Interleaving



PennState

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Avoiding race conditions



PennState

- To avoid race conditions we need to ensure that only a single thread can execute a critical section at any given time
- For simple cases we can use atomics (`#include <stdatomic.h>`)
 - modifying a variable res
- In general, however, a complex logic
- We need primitives for mutual exclusion that only one thread is executing the critical section while others are prevented from doing so
- One way to achieve mutual exclusion is using locks:

```
lock_t mutex
lock(&mutex)
critical section
unlock(&mutex)
```

Threads accessing shared data



PennState

- Fixing race condition using atomics

```
#include <stdatomic.h>
```

```
static atomic_int counter = 0;
```

```
void *func(void *arg)
```

```
    for (int i = 0; i < 10; i++)  
        ++counter;
```

```
    return NULL;
```

```
}
```

```
int main(void)
```

```
{
```

```
    pthread_t t1, t2;
```

```
    printf("counter = %d\n", counter);
```

```
    pthread_create(&t1, NULL, func, NULL);
```

```
    pthread_create(&t2, NULL, func, NULL);
```

```
    pthread_join(t1, NULL);
```

```
    pthread_join(t2, NULL);
```

```
    printf("both threads completed, counter = %d\n", counter);
```

```
    return 0;
```

```
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Threads accessing shared data – Fixed!



PennState

- Fixing race condition using mutexes

```
static int counter = 0;
pthread_mutex_lock_t lock = PTHREAD_MUTEX_INITIALIZER;

void *func(void *arg) {
    for (int i = 0; i < 5000; ++i) {
        pthread_mutex_lock
        ++counter;
        pthread_mutex_unlo
    }
    return NULL;
}

int main(void)
{
    pthread_t t1, t2;
    printf("counter = %d\n", counter);
    pthread_create(&t1, NULL, func, NULL);
    pthread_create(&t2, NULL, func, NULL);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("both threads completed, counter = %d\n", counter);
    return 0;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Threads tradeoffs



PennState

- Benefits

- still the case that much of the code is identical!
- parallel execution; good CPU, network utilization
- lower overhead than pr
- shared-memory commu

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Disadvantages

- synchronization is complicated
- shared fate within a process; one rogue thread can hurt you
- security (no isolation)

Add WeChat edu_assist_pro

- We scratched the surface – more advanced usage will be taught in 473