



PennState

# CMPSC 311 - Introduction to Systems Programming

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Introduction to Profiling  
Add WeChat edu\_assist\_pro

Suman Saha

(Slides are mostly by Professors Patrick McDaniel and  
Abutalib Aghayev)

---

# Program Performance



PennState

- Programs run only as well as the code you write
- Poor code often runs poorly
  - Crashes or generates incorrect output (bugs)
  - Is laggy, jittery or slow (
    - Too slow on real inputs (data)
    - Not-reactive enough to be usable (interface)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Optimization



PennState

- Optimization is the process where you take an existing program and alter it to remove inefficiencies.

- Change algorithms
- Restructure code
- Redesign data structures
- Refactor code

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

## Career Notes:

1. Learning to optimize your code is essential to becoming a professional programmer.
2. Optimizing code is a phase of development you don't experience in school.
3. You will spend a good deal of your professional life optimizing existing code without changing its function.

# Don't optimize prematurely



PennState

- The first rule of optimization is:

- Don't do it.

- The second rule of optimization (for experts only):

- Don't do it yet. Measure

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- It is far, far easier to make a correct program  
than it is to make a fast program correct

# Premature optimization is the root of all..



PennState

- "Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at have a strong negative debugging and maintenance are consid We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil (quoting Tony Hoare). Yet we should not pass up our opportunities in that critical 3%."

Donald Knuth

# Example Inefficient Code



PennState

```
uint64_t prod_one(uint64_t a, uint64_t b) {
    uint64_t out = 0;
    for (uint64_t i = 0; i < a; ++i)
        out += b;
    return out;
}
```

```
uint64_t prod_two(uint64_t a, uint64_t b) {
    return a * b;
}
```

```
int main(void) {
    uint64_t a = 10000000000, b = 100000000, sum = 0;
    clock_t start;
    double elapsed;

    start = clock();
    sum = prod_one(a, b);
    elapsed = (clock() - start) / CLOCKS_PER_SEC;
    printf("%ld * %ld = %ld (took %.3f seconds) \n", a, b, sum, elapsed);

    start = clock();
    sum = prod_two(a, b);
    elapsed = (clock() - start) / CLOCKS_PER_SEC;
    printf("%ld * %ld = %ld (took %.3f seconds) \n", a, b, sum, elapsed);

    return 0;
}
```

- Try it with
  - gcc -O0
  - gcc -O1
  - gcc -O2

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

```
$ gcc -O2 -Wall opt.c -o opt
[0s euclid:~/tmp (master)]
$ ./opt
10000000000 + 100000000 = 1000000000000000000 (took 0.000 seconds)
10000000000 + 100000000 = 1000000000000000000 (took 0.000 seconds)
[0s euclid:~/tmp (master)]
```

# Profiling



PennState

- Debugging helps the programmer find and fix bugs ...
- **Profiling** helps the programmer find inefficiencies
  - Profiling involves running a version of the program instrumented with code to measure how much time is spent
  - How much time in each module of the program?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



- **gprof** is a utility that measures a program's performance and behavior.
- This produces **non-interactive** profile output
- The output provides detail
  - The time that the program took and time for each function
    - Statistics and detail on "performance"
  - Which functions called other functions and how many times
    - Statistics and detail on the "call graph"

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



# Running gprof



PennState

1. First compile program using the “-pg” flag:

```
$ gcc -pg profiling.c -o profiling
```

2. Run the program (will generate file gmon.out):

```
$ ./pr
```

3. Run gprof with the command <https://eduassistpro.github.io/>

```
$ gprof profiling |
```

4. Review the output [Add WeChat edu\\_assist\\_pro](#)

```
$ gprof profiling
Flat profile:
...
```

5. Optimize the program, re-profile
6. GOTO step#1

# Gprof (flat profile)



PennState

Assignment Project Exam Help

```
$ gprof opt
Flat profile:
```

Each sample counts as

% time	cumulative seconds	self seconds				
101.31	21.87	21.87	1	21.87	21.87	prod_one
0.00	21.87	0.00	1	0.00	0.00	

%  
time      the percentage of the total running time of  
          program used by this function.

cumulative a running sum of the number of seconds accounted  
seconds    for by this function and those listed above it.

self       the number of seconds accounted for by this  
seconds    function alone. This is the major sort for this  
          listing.

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Gprof (Call Graph)



PennState

Assignment Project Exam Help

Call graph (explanation follows)

granularity: each sample

index	% time	self	ch		
		21.87	0.00	1/1	main [2]
[1]	100.0	21.87	0.00	1	prod_one
-----					
[2]	100.0	0.00	21.87		main [2]
		21.87	0.00	1/1	prod_one [1]
		0.00	0.00	1/1	prod_two [3]
-----					
		0.00	0.00	1/1	main [2]
[3]	0.0	0.00	0.00	1	prod_two [3]
-----					

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro


# Optimization revisited ...



PennState

- When optimizing, you focus on modules of the program which implement the features and processing of the program.

- Which parts of the program depends on what parts
- then, focus on those parts which take up the most time.



- Profiling tells us where to spend our time.

# Amdahl's Law



PennState

- Amdahl's law models the maximum performance gain that can be expected by improving part of the system, i.e., what can we expect in terms of improvement.

Assignment Project Exam Help

- Consider

- $k$  is the percentage of  $t$  optimized module(s).
- $s$  is the execution time expressed in terms speedup (2X, 3X...), which can be found as

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

$$s = \frac{\text{original execution time}}{\text{improved execution time}}$$

# Amdahl's Law (cont.)



- The overall speedup  $T$  of the program is expressed :

$$T = \frac{1}{(1 - \frac{k}{s}) + \frac{k}{s}}$$

Assignment Project Exam Help

- Intuition:

- $1 - \frac{k}{s}$  is the part of the

<https://eduassistpro.github.io/>

- $\frac{k}{s}$  is the ratio of altered program size to sp

Add WeChat edu\_assist\_pro

# Amdahl's Law (example)



- Assume that a module A of a program is optimized.

- A represents 45% of the run time of the program.
- The optimization reduces the runtime of module from 750ms to 50ms.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- What is the program speedup?

# Amdahl's Law (example)



PennState

- Assume that a module A of a program is optimized.

- A represents 45% of the run time of the program.
- The optimization reduces the runtime of module from 750ms to 50ms.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

$$T = \frac{1}{(1-.45) + \frac{.45}{15}} = \frac{1}{.55 + .03} = 1.724$$

- What is the program speedup? (A: 1.724X)



# A more complex example



PennState

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# A even more complex example:



- Assume another system: we have 4 modules each being measured before and after optimization

Module	Before Optimization (usec)	After Optimization (usec)
		60
		11
		600
D		1

- Now suppose that the runtime of the original execution is 2000 usec, what is the speedup?

$$T = \frac{1}{(1 - k) + \frac{k}{s}}$$

# What is going on?



PennState

```
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>

#define SIZE 35000

int main(void) {
    uint64_t **a = malloc(SIZE * sizeof(uint64_t *));
    for (int i = 0; i < SIZE; ++i)
        a[i] = malloc(SIZE * sizeof(uint64_t));

    uint64_t sum = 0;
    for (int i = 0; i < SIZE; ++i)
        for (int j = 0; j < SIZE; ++j)
            sum += a[i][j];

    printf("sum is %ld\n", sum);

    return 0;
}
```

```
$ gcc -Wall -O2 c.c -o c
$ time ./c
sum is 0
```

```
real    0m1.560s
user    0m1.060s
sys     0m0.500s
```

```
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>

#define SIZE 35000

int main(void) {
    uint64_t **a = malloc(SIZE * sizeof(uint64_t *));
    for (int i = 0; i < SIZE; ++i)
        a[i] = malloc(SIZE * sizeof(uint64_t));

    uint64_t sum = 0;
    for (int i = 0; i < SIZE; ++i)
        for (int j = 0; j < SIZE; ++j)
            sum += a[i][j];

    printf("sum is %ld\n", sum);

    return 0;
}
```

```
$ gcc -Wall -O2 c.c -o c
$ time ./c
sum is 0
```

```
real    0m14.543s
user    0m13.597s
sys     0m0.948s
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro